

# Week 1 Reflections

Prem Kumar Mani

In this exercise, I tried to create four agents and their corresponding tasks. Initially, I had some confusion regarding the goal and backstory of each agent. I started by defining one agent, followed by its corresponding task. I then ran the script, ensured that the agent worked as expected, and moved on to the next agent.

The agents produced the desired output. I then tried different queries that were not present in, or were irrelevant to, the available documents and noticed that the agents still produced results. This showed me that unless an agent's goal is clear and specific, it may return information from other sources or from its own general knowledge, which may not be relevant to the intended task.

I had to refine the goal for the **source\_hunter** agent to strictly search the available documents and return results only when a relevant document reference was available. I also initially used **memory=True** for all agents without a specific reason. I realized that memory is not particularly relevant for this workflow because every search query can be treated independently of previous search queries.

Regarding the output, I had to provide a separate instruction to write the final report to an output file and refine the filename format. I am still working on getting the filename exactly as I want it.

I also made the research question dynamic so that users can provide different questions. However, I did not implement many guardrails to prevent irrelevant or inappropriate input questions.

Overall, this exercise gave me better clarity on defining boundaries for each agent and assigning separate roles and objectives. I also learned how the output of one agent can become the input for the next agent in a multi-agent workflow.

To make this solution production-ready, I would introduce stronger guardrails, define clearer roles and goals for each agent, enforce stricter document grounding, and optimize the time taken to generate the final report.